

HW6 – Part 2: Neural Network on Image Classification

Nafis Neehal

661990881

Code Link in Colab:

<https://colab.research.google.com/drive/1-S75WAemnDE-knabrNqouwwWgVz6qjcl>

Data Description

Table 1: Data Description

Total training data provided	4000
Total test data provided	1000
Validation to Training Ratio	12.5%
Final number of Validation Data	500
Final number of Training Data	3500
Mini Batch Size (for train, validation, test)	64
Data shape	64*3*60*90 (aspect ratio 1.5 = height/width)

- ✓ I would have gone for the standard 60-20-20 (3k-1k-1k for this dataset) split for train-validation-test. But the number of training examples are inadequate to gain a good accuracy without the network overfitting with only 3k training examples. So, I tried to maximize the number of training examples as much as I could and didn't want to manipulate the test set at all. It might have hampered the validation legitimacy a bit, but at least the network is not that overfitted.
- ✓ I explored data-augmentation for increasing the number of training examples to reduce the effect of overfitting and also introducing the 60-20-20 split, but as I don't have any GPU in my local machine, it was a bit tough to do that using Google Colab. I tried doing it in runtime using "torchvision.transforms" and tried just introducing Horizontal and Vertical Flips, but it ended up consuming a huge amount of time.
- ✓ I applied the following transformations on each data –
 - Resize each image to (3*60*90) maintain aspect ratio of 1.5
 - Converted each image array to tensors
 - Normalize the images
- ✓ I downsampled the data 2 times (one-fourth the size) and got the best performance that I could I achieve from both FCNN and CNN.
- ✓ I normalized the data by setting the mean to [0.5, 0.5, 0.5] and std to [0.5, 0.5, 0.5]. I found this value works best in several articles and literature.
- ✓ I used "torch.randperm" to shuffle all the data before assigning it to training/validation set with seed value = 1. Maybe trying some different seed values would have been another parameter search task, but I just don't have the computation power to do so, and so couldn't tweak and analyze different seed values.
- ✓ I used torchvision.datasets.ImageFolder to load my custom image dataset from Google Drive.

- ✓ I used `torchvision.utils.data.DataLoader` with `num_workers = 4` for parallel data transfer with the setting `pin_memory=True` which also helped with faster data transfer between CPU and GPU (suggestion from NVidia developer blog and PyTorch forum)
- ✓ I loaded all my data from Google drive after I mounted the drive in Colab.

Network Common Parameters

Table 2: Common Network Parameters

Number of classes	5
Model Name	NN / CNN
Number of Epochs	50
Learning Rate	1e-3
Device	GPU (if available) / CPU
Optimizer	Adam
L2 Regularization with Weight Decay	1e-5
Loss Function	Cross Entropy Loss

Overall issues and general brief process of Train, Validation and Test:

- ✓ In each epoch, I trained the data with all the mini-batches (size = 64), then validate the performance on the validation data. Validation could have been done periodically (i.e., after each 5 epoch), but I did in each epoch to track my model's performance more accurately, specially, it helped tackling the overfitting.
- ✓ I initially cleared all the gradients after the loss was propagated backwards which literally skyrocketed the gradients. Later I fixed it by calling the `optimizer.zero_grad()` before loss was propagated.
- ✓ No gradient update or loss propagate in validation / test; while `model.eval()` is on.
- ✓ Pushed all data tensors and my network model to GPU. Used Tesla K-80 (Google Colab's default).
- ✓ In each epoch, kept the track of the best model. If my models validation accuracy increased and validation loss decreased than the current best saved values(highest validation accuracy and lowest loss), I saved it as best model using `torch.save()` with all parameter states of the model and the optimizer. Then while running final test on the test dataset, I loaded that best model and use that to evaluate the final performance of my network on test data. The epoch value of the final best model state can be used for the threshold for early stopping in order to overfit the network.

General Measures to reduce overfitting

- ✓ L2 regularization with weight decay 1e-5
- ✓ Dropout ($p = 0.5$)
- ✓ Tried Early stopping as well, but this did not fit well as the model tend to overfit very early (only after 5-6 epochs)
- ✓ Tried to reduce overfitting also by reducing the model complexity (simpler architecture)

HW6 – Part 2(a): Fully Connected Neural Network on Image Classification

Network Architecture:

This image is a rough representation of the network. Exact number of nodes were not possible to show here.




Fig 1: FCNN Representation

Network Description:

Table 3: FCNN Network Architecture

Layer	Activation	Number of Weights	Number of Biases	Dimension
Input	None	0	0	1x16200 (3*60*90)
Hidden 1	ReLU	16200*512	512	1x512
Hidden 2	ReLU	512*256	256	1x256
Hidden 3	ReLU	256*128	128	1x128
Output	LogSoftMax	128*5	5	1x5

- ✓ I used dropout with $P=0.5$ between all the layers (H1-H2, H2-H3, H3-Output). Helped me with the overfitting. I tried reducing one hidden layer to make the architecture simpler, but then the validation and test accuracy gets severely affected.
- ✓ I explored LeakyReLU activation too for hidden layers, but ReLU performed better.
- ✓ I used CrossEntropyLoss for loss function, and torch.nn.CrossEntropyLoss already contains nn.LogSoftMax which was by default applied on the output of the output layer, that is why I explicitly didn't use any other SoftMax.

Performance Analysis:




Fig 2: Training and Validation Data in Graph for FCNN

Epoch[1/50]	Training Accuracy: 0.5300	Training Loss 0.0192	Validation Accuracy: 0.7120	Validation Loss: 0.0128
Epoch[2/50]	Training Accuracy: 0.6677	Training Loss 0.0148	Validation Accuracy: 0.7320	Validation Loss: 0.0120
Epoch[3/50]	Training Accuracy: 0.7066	Training Loss 0.0126	Validation Accuracy: 0.7380	Validation Loss: 0.0114
Epoch[4/50]	Training Accuracy: 0.7220	Training Loss 0.0121	Validation Accuracy: 0.7620	Validation Loss: 0.0118
Epoch[5/50]	Training Accuracy: 0.7434	Training Loss 0.0116	Validation Accuracy: 0.7520	Validation Loss: 0.0112
Epoch[6/50]	Training Accuracy: 0.7526	Training Loss 0.0111	Validation Accuracy: 0.7460	Validation Loss: 0.0111
Epoch[7/50]	Training Accuracy: 0.7623	Training Loss 0.0105	Validation Accuracy: 0.7460	Validation Loss: 0.0111
Epoch[8/50]	Training Accuracy: 0.7737	Training Loss 0.0099	Validation Accuracy: 0.7480	Validation Loss: 0.0114
Epoch[9/50]	Training Accuracy: 0.7857	Training Loss 0.0099	Validation Accuracy: 0.7500	Validation Loss: 0.0113
Epoch[10/50]	Training Accuracy: 0.7837	Training Loss 0.0094	Validation Accuracy: 0.7560	Validation Loss: 0.0116
Epoch[11/50]	Training Accuracy: 0.7894	Training Loss 0.0093	Validation Accuracy: 0.7540	Validation Loss: 0.0112
Epoch[12/50]	Training Accuracy: 0.8063	Training Loss 0.0088	Validation Accuracy: 0.7540	Validation Loss: 0.0114
Epoch[13/50]	Training Accuracy: 0.8097	Training Loss 0.0087	Validation Accuracy: 0.7460	Validation Loss: 0.0119
Epoch[14/50]	Training Accuracy: 0.8086	Training Loss 0.0086	Validation Accuracy: 0.7500	Validation Loss: 0.0114
Epoch[15/50]	Training Accuracy: 0.8169	Training Loss 0.0083	Validation Accuracy: 0.7560	Validation Loss: 0.0111
Epoch[16/50]	Training Accuracy: 0.8157	Training Loss 0.0083	Validation Accuracy: 0.7220	Validation Loss: 0.0121
Epoch[17/50]	Training Accuracy: 0.8243	Training Loss 0.0080	Validation Accuracy: 0.7440	Validation Loss: 0.0122
Epoch[18/50]	Training Accuracy: 0.8166	Training Loss 0.0078	Validation Accuracy: 0.7620	Validation Loss: 0.0127
Epoch[19/50]	Training Accuracy: 0.8423	Training Loss 0.0074	Validation Accuracy: 0.7600	Validation Loss: 0.0125
Epoch[20/50]	Training Accuracy: 0.8409	Training Loss 0.0071	Validation Accuracy: 0.7380	Validation Loss: 0.0132
Epoch[21/50]	Training Accuracy: 0.8391	Training Loss 0.0072	Validation Accuracy: 0.7500	Validation Loss: 0.0128
Epoch[22/50]	Training Accuracy: 0.8451	Training Loss 0.0071	Validation Accuracy: 0.7540	Validation Loss: 0.0129
Epoch[23/50]	Training Accuracy: 0.8474	Training Loss 0.0066	Validation Accuracy: 0.7460	Validation Loss: 0.0134
Epoch[24/50]	Training Accuracy: 0.8486	Training Loss 0.0068	Validation Accuracy: 0.7560	Validation Loss: 0.0129
Epoch[25/50]	Training Accuracy: 0.8557	Training Loss 0.0066	Validation Accuracy: 0.7560	Validation Loss: 0.0127
Epoch[26/50]	Training Accuracy: 0.8549	Training Loss 0.0065	Validation Accuracy: 0.7320	Validation Loss: 0.0129
Epoch[27/50]	Training Accuracy: 0.8537	Training Loss 0.0062	Validation Accuracy: 0.7660	Validation Loss: 0.0138
Epoch[28/50]	Training Accuracy: 0.8591	Training Loss 0.0064	Validation Accuracy: 0.7640	Validation Loss: 0.0138
Epoch[29/50]	Training Accuracy: 0.8531	Training Loss 0.0064	Validation Accuracy: 0.7560	Validation Loss: 0.0136
Epoch[30/50]	Training Accuracy: 0.8654	Training Loss 0.0061	Validation Accuracy: 0.7500	Validation Loss: 0.0142
Epoch[31/50]	Training Accuracy: 0.8540	Training Loss 0.0062	Validation Accuracy: 0.7760	Validation Loss: 0.0158
Epoch[32/50]	Training Accuracy: 0.8680	Training Loss 0.0059	Validation Accuracy: 0.7540	Validation Loss: 0.0149
Epoch[33/50]	Training Accuracy: 0.8560	Training Loss 0.0064	Validation Accuracy: 0.7600	Validation Loss: 0.0147
Epoch[34/50]	Training Accuracy: 0.8683	Training Loss 0.0062	Validation Accuracy: 0.7540	Validation Loss: 0.0147
Epoch[35/50]	Training Accuracy: 0.8706	Training Loss 0.0059	Validation Accuracy: 0.7720	Validation Loss: 0.0137
Epoch[36/50]	Training Accuracy: 0.8751	Training Loss 0.0057	Validation Accuracy: 0.7520	Validation Loss: 0.0140
Epoch[37/50]	Training Accuracy: 0.8791	Training Loss 0.0054	Validation Accuracy: 0.7640	Validation Loss: 0.0144
Epoch[38/50]	Training Accuracy: 0.8846	Training Loss 0.0057	Validation Accuracy: 0.7540	Validation Loss: 0.0138
Epoch[39/50]	Training Accuracy: 0.8809	Training Loss 0.0057	Validation Accuracy: 0.7580	Validation Loss: 0.0138
Epoch[40/50]	Training Accuracy: 0.8797	Training Loss 0.0053	Validation Accuracy: 0.7500	Validation Loss: 0.0144
Epoch[41/50]	Training Accuracy: 0.8886	Training Loss 0.0053	Validation Accuracy: 0.7480	Validation Loss: 0.0142
Epoch[42/50]	Training Accuracy: 0.8920	Training Loss 0.0051	Validation Accuracy: 0.7480	Validation Loss: 0.0150
Epoch[43/50]	Training Accuracy: 0.8946	Training Loss 0.0048	Validation Accuracy: 0.7720	Validation Loss: 0.0156
Epoch[44/50]	Training Accuracy: 0.8977	Training Loss 0.0045	Validation Accuracy: 0.7540	Validation Loss: 0.0168
Epoch[45/50]	Training Accuracy: 0.8929	Training Loss 0.0051	Validation Accuracy: 0.7500	Validation Loss: 0.0159
Epoch[46/50]	Training Accuracy: 0.8837	Training Loss 0.0051	Validation Accuracy: 0.7560	Validation Loss: 0.0168
Epoch[47/50]	Training Accuracy: 0.8820	Training Loss 0.0055	Validation Accuracy: 0.7540	Validation Loss: 0.0154
Epoch[48/50]	Training Accuracy: 0.8943	Training Loss 0.0049	Validation Accuracy: 0.7540	Validation Loss: 0.0155
Epoch[49/50]	Training Accuracy: 0.8983	Training Loss 0.0046	Validation Accuracy: 0.7480	Validation Loss: 0.0166
Epoch[50/50]	Training Accuracy: 0.9046	Training Loss 0.0046	Validation Accuracy: 0.7640	Validation Loss: 0.0157

Fig 3: Training and Validation Log for FCNN

From the very first sight, it seems obvious that the network is overfitted, at least that's what the graph says. Validation accuracy almost flattens after a certain number of epochs and the trend of the graph of validation loss is going upwards. But if we closely observe the rate of increase in validation loss, it is very negligible.

Best result –

Table 4: Final FCNN Performance

Epoch No	15
Training Accuracy	81.685%
Train Loss	0.0083
Validation Accuracy	75.60%
Validation Loss	0.0111
Test Accuracy	77.70%

Table 5: Classwise Accuracy

Class Name	Accuracy on Test Data
Grass	76%
Ocean	58%
Red carpet	92%
Road	75%
Wheatfield	70%

The reason that the train accuracy is lower than typical rate (above 95% in overfitted networks) is because of the dropout in between the FC layers. Initially, when I trained and validated the network for the first time, the train accuracy was around 98% and validation was around 65% and the validation loss was skyrocketing while train loss was going downwards! So, by tuning the parameters in order to trying and balance the network overfitting, I was able to increase the result up to around 12%.

Although, from the result using the model I chose as best is not that impressive, but the result is at least balanced and states the fact that the model is not overfitted. Epoch No 15 represents here that this is the threshold for “Early Stopping” after which network starts to perform badly. After this point, the validation loss starts to go upwards and validation accuracy flattens with spikes (accuracy fluctuation because of using dropout).




Fig 4: Training and Validation Log for FCNN up to early stopping threshold (Epoch 15)

HW6 – Part 2(b): Convolutional Neural Network on Image Classification

Network Architecture:

This image is a rough representation of the network. Exact number of nodes were not possible to show here.




Fig 5: Training and Validation Log for FCNN up to early stopping threshold (Epoch 15)

Network Description:

Table 6: CNN Network Description

Layer	Unit	Input Shape (batch_size*depth*height*width)	Output Shape (batch_size*depth*height*width)	Filter size
Convolution	1	32x3x60x90	32x32x60x90	3x3
Max-Pool		32x32x60x90	32x32x30x45	2x2
Convolution	2	32x32x30x45	32x64x30x45	3x3
Max-Pool		32x64x30x45	32x64x15x22	2x2
Convolution	3	32x64x15x22	32x128x15x22	3x3
Max-Pool		32x128x15x22	32x128x7x11	2x2
Dense-1	4	32x128x7x11	32x1x256x1	N/A
Dense-2	5	32x1x256x1	32x1x5x1	N/A

I have also used –

- Dropout(p=0.5) before each of the Dense layers. I explored with p=0.2 and p=0.8, but p=0.5 gave the optimum result
- There is batch-normalization after each Convolution layer in each of the first 3 units (ConvUnits).
- I used 3x3 filters for Convolution and 2x2 filters for Max-Pool. I would have explored more with different filter size but had limitation in computational resource.
- ReLU was applied after each BatchNorm was applied.

Performance Analysis:




Fig 6: Training and Validation Data in Graph for CNN

Epoch[1/50]	Training Accuracy: 0.6557	Training Loss 0.0345	Validation Accuracy: 0.7740	Validation Loss: 0.0166
Epoch[2/50]	Training Accuracy: 0.7649	Training Loss 0.0139	Validation Accuracy: 0.7800	Validation Loss: 0.0139
Epoch[3/50]	Training Accuracy: 0.7946	Training Loss 0.0102	Validation Accuracy: 0.7880	Validation Loss: 0.0120
Epoch[4/50]	Training Accuracy: 0.8183	Training Loss 0.0089	Validation Accuracy: 0.7840	Validation Loss: 0.0104
Epoch[5/50]	Training Accuracy: 0.8323	Training Loss 0.0079	Validation Accuracy: 0.7720	Validation Loss: 0.0114
Epoch[6/50]	Training Accuracy: 0.8469	Training Loss 0.0069	Validation Accuracy: 0.8280	Validation Loss: 0.0090
Epoch[7/50]	Training Accuracy: 0.8517	Training Loss 0.0069	Validation Accuracy: 0.8080	Validation Loss: 0.0091
Epoch[8/50]	Training Accuracy: 0.8569	Training Loss 0.0067	Validation Accuracy: 0.8200	Validation Loss: 0.0109
Epoch[9/50]	Training Accuracy: 0.8709	Training Loss 0.0061	Validation Accuracy: 0.7800	Validation Loss: 0.0117
Epoch[10/50]	Training Accuracy: 0.8743	Training Loss 0.0058	Validation Accuracy: 0.8280	Validation Loss: 0.0095
Epoch[11/50]	Training Accuracy: 0.8809	Training Loss 0.0056	Validation Accuracy: 0.8180	Validation Loss: 0.0096
Epoch[12/50]	Training Accuracy: 0.8777	Training Loss 0.0056	Validation Accuracy: 0.8120	Validation Loss: 0.0088
Epoch[13/50]	Training Accuracy: 0.8891	Training Loss 0.0049	Validation Accuracy: 0.8060	Validation Loss: 0.0111
Epoch[14/50]	Training Accuracy: 0.8814	Training Loss 0.0057	Validation Accuracy: 0.8280	Validation Loss: 0.0095
Epoch[15/50]	Training Accuracy: 0.8880	Training Loss 0.0048	Validation Accuracy: 0.8300	Validation Loss: 0.0094
Epoch[16/50]	Training Accuracy: 0.8994	Training Loss 0.0044	Validation Accuracy: 0.8380	Validation Loss: 0.0085
Epoch[17/50]	Training Accuracy: 0.8949	Training Loss 0.0045	Validation Accuracy: 0.8220	Validation Loss: 0.0096
Epoch[18/50]	Training Accuracy: 0.9086	Training Loss 0.0043	Validation Accuracy: 0.8360	Validation Loss: 0.0091
Epoch[19/50]	Training Accuracy: 0.9046	Training Loss 0.0045	Validation Accuracy: 0.8280	Validation Loss: 0.0095
Epoch[20/50]	Training Accuracy: 0.9003	Training Loss 0.0047	Validation Accuracy: 0.8120	Validation Loss: 0.0107
Epoch[21/50]	Training Accuracy: 0.9083	Training Loss 0.0042	Validation Accuracy: 0.8360	Validation Loss: 0.0098
Epoch[22/50]	Training Accuracy: 0.9171	Training Loss 0.0038	Validation Accuracy: 0.8440	Validation Loss: 0.0104
Epoch[23/50]	Training Accuracy: 0.9106	Training Loss 0.0039	Validation Accuracy: 0.8160	Validation Loss: 0.0110
Epoch[24/50]	Training Accuracy: 0.9103	Training Loss 0.0039	Validation Accuracy: 0.8220	Validation Loss: 0.0110
Epoch[25/50]	Training Accuracy: 0.9137	Training Loss 0.0038	Validation Accuracy: 0.8200	Validation Loss: 0.0102
Epoch[26/50]	Training Accuracy: 0.9217	Training Loss 0.0040	Validation Accuracy: 0.8280	Validation Loss: 0.0097
Epoch[27/50]	Training Accuracy: 0.9297	Training Loss 0.0033	Validation Accuracy: 0.8300	Validation Loss: 0.0104
Epoch[28/50]	Training Accuracy: 0.9186	Training Loss 0.0036	Validation Accuracy: 0.8200	Validation Loss: 0.0135
Epoch[29/50]	Training Accuracy: 0.9423	Training Loss 0.0027	Validation Accuracy: 0.8320	Validation Loss: 0.0114
Epoch[30/50]	Training Accuracy: 0.9474	Training Loss 0.0026	Validation Accuracy: 0.8200	Validation Loss: 0.0141
Epoch[31/50]	Training Accuracy: 0.9466	Training Loss 0.0025	Validation Accuracy: 0.7940	Validation Loss: 0.0151
Epoch[32/50]	Training Accuracy: 0.9300	Training Loss 0.0029	Validation Accuracy: 0.8600	Validation Loss: 0.0112
Epoch[33/50]	Training Accuracy: 0.9503	Training Loss 0.0021	Validation Accuracy: 0.7940	Validation Loss: 0.0142
Epoch[34/50]	Training Accuracy: 0.9423	Training Loss 0.0025	Validation Accuracy: 0.8380	Validation Loss: 0.0123
Epoch[35/50]	Training Accuracy: 0.9483	Training Loss 0.0025	Validation Accuracy: 0.8320	Validation Loss: 0.0128
Epoch[36/50]	Training Accuracy: 0.9543	Training Loss 0.0021	Validation Accuracy: 0.8340	Validation Loss: 0.0128
Epoch[37/50]	Training Accuracy: 0.9531	Training Loss 0.0022	Validation Accuracy: 0.8500	Validation Loss: 0.0128
Epoch[38/50]	Training Accuracy: 0.9480	Training Loss 0.0022	Validation Accuracy: 0.8220	Validation Loss: 0.0163
Epoch[39/50]	Training Accuracy: 0.9563	Training Loss 0.0019	Validation Accuracy: 0.8260	Validation Loss: 0.0143
Epoch[40/50]	Training Accuracy: 0.9609	Training Loss 0.0017	Validation Accuracy: 0.8280	Validation Loss: 0.0131
Epoch[41/50]	Training Accuracy: 0.9660	Training Loss 0.0019	Validation Accuracy: 0.8440	Validation Loss: 0.0135
Epoch[42/50]	Training Accuracy: 0.9737	Training Loss 0.0012	Validation Accuracy: 0.8440	Validation Loss: 0.0128
Epoch[43/50]	Training Accuracy: 0.9754	Training Loss 0.0011	Validation Accuracy: 0.8420	Validation Loss: 0.0160
Epoch[44/50]	Training Accuracy: 0.9689	Training Loss 0.0013	Validation Accuracy: 0.8460	Validation Loss: 0.0132
Epoch[45/50]	Training Accuracy: 0.9689	Training Loss 0.0013	Validation Accuracy: 0.8280	Validation Loss: 0.0154
Epoch[46/50]	Training Accuracy: 0.9669	Training Loss 0.0015	Validation Accuracy: 0.8320	Validation Loss: 0.0144
Epoch[47/50]	Training Accuracy: 0.9703	Training Loss 0.0014	Validation Accuracy: 0.8440	Validation Loss: 0.0162
Epoch[48/50]	Training Accuracy: 0.9646	Training Loss 0.0017	Validation Accuracy: 0.8180	Validation Loss: 0.0197
Epoch[49/50]	Training Accuracy: 0.9597	Training Loss 0.0019	Validation Accuracy: 0.8380	Validation Loss: 0.0139
Epoch[50/50]	Training Accuracy: 0.9694	Training Loss 0.0016	Validation Accuracy: 0.8360	Validation Loss: 0.0150

Fig 7: Training and Validation Log for CNN

For CNN, although the training accuracy goes all the way up to 97% after full training(50 epochs), I chose to early stop at epoch 16, after which the validation loss starts the generic trend of going up in each epoch while the training loss goes downwards which is the prime sign of overfitting. The reason for choosing a version of the network that is comparatively performing less accurately than the final model is to avoid overfitting; which means the version I chose as best would have better performance in generalizing any unknown data than the final one.

Table 7: Final CNN Performance

Epoch	16
Training Accuracy	89.943%
Training Loss	0.004
Validation Accuracy	83.800%
Validation Loss	0.008
Test Accuracy	85.900%

Table 8: Classwise Accuracy:

Class Name	Accuracy on Test Data
Grass	64%
Ocean	100%
Red carpet	91%
Road	100%
Wheatfield	82%




Fig 8: Training and Validation Log for CNN up to early stopping threshold (Epoch 16)

Concluding Remarks:

I could have tweaked the network better if I had the computational resources. Google Colab doesn't allow unlimited usage of GPU and I got banned 2-3 times for 12 hour period as I have been continuously performing arduous computations for the last 1 week.