


CSci 4270 and 6270  
Spring 2021  
Lecture 7: Feature Extraction — Edges  
Thursday, February 19, 2021

## Overview

- Goal: extract image locations of sharpest intensity change.
  - Intuitively, edges are “significant events” in images.
  - This is a “boundary first” approach to image analysis, whereas the problem area known as *segmentation* is a “region first” approach.
- A single pixel edge is called an *edgel*, short for “edge element”.
- Steps of edge(l) detection
  - Smoothing
  - Differentiation, gradient magnitude and direction
    - \* Some methods use second order derivatives, but we will not.
  - Non-maximum suppression, sub-pixel localization, thresholding
  - Linking into edgel chains
- We will discuss each in turn.



## Smoothing for Edge Detection

Choice is the Gaussian, for reasons we have already outlined

- Isotropic and separable
- Fast implementations
- Good spatial and frequency properties

## Choosing the Smoothing Scale: Scale Space

- Smaller values of  $\sigma$  lead to more detailed edges and better edge localization.
- Larger values of  $\sigma$  lead to fewer and perhaps more significant edges
- Different scales are needed when images have different sizes or magnifications
- We can combine these by working with multiple scales, treating  $\sigma$  as an additional variable, producing what is called *scale space*.
  - We will look at scale space in much more detail in future lectures.
- For the purposes of these notes, we will work with the results of smoothing with a single value of  $\sigma$ .

## Differential Operations

- Two common choices
  - 1st derivative based on the directional derivatives, the gradient, and finding peak.
  - 2nd derivative based on the Laplacian and finding “zero-crossings”.
- While our emphasis will be on first derivative operations, during lecture we will take a brief look at the formal structure of the Laplacian and why it is used. In particular, we will discuss the equation


$$\nabla^2(I * G)$$

where

$$\nabla^2(f) = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

*im? ↗ gaussian smootn*  
*compute 2nd derivative and add result*

We will return to the use of the Laplacian during “interest point” extraction.



## Computing the Gradient

- Let  $I_s$  be the smoothed image.
- Compute  $I_x = \partial I_s / \partial x$  and  $I_y = \partial I_s / \partial y$  using discrete differentiation.
- Then we can compute the gradient direction and magnitude as

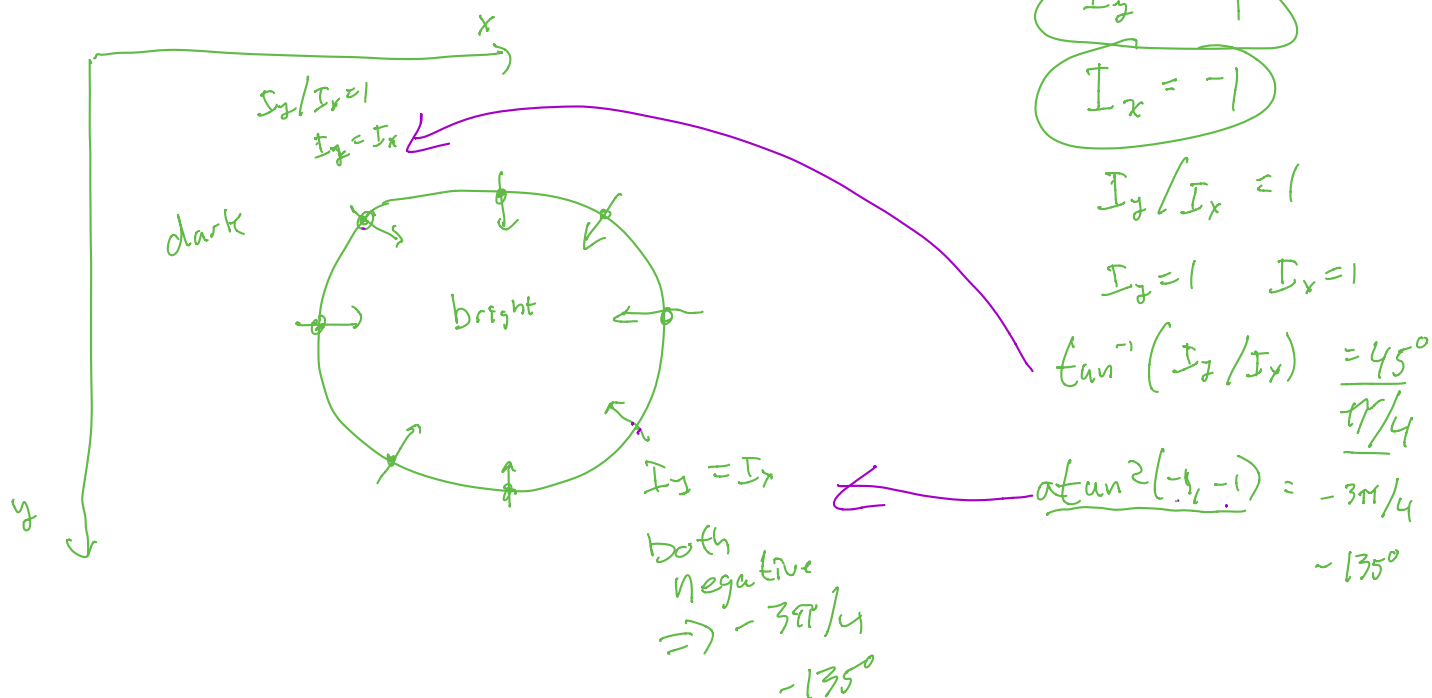
$$I_m(x, y) = \|I_x(x, y)^2 + I_y(x, y)^2\|^{1/2}$$

$$I_\theta(x, y) = \text{atan2}(I_y(x, y), I_x(x, y))$$

range  $[-\pi, \pi)$

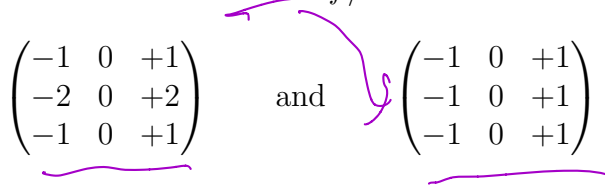
– Notice with atan2 the y derivative is first.

- The direction is “normal” to the direction of the edge contour and points in the direction of most rapid intensity increase.



## Other Operators

- As we've discussed, other partial derivative operators have been used. For example, the Sobel and Prewitt  $\partial f / \partial x$  kernels are




$$\begin{pmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{pmatrix} \quad \text{and} \quad \begin{pmatrix} -1 & 0 & +1 \\ -1 & 0 & +1 \\ -1 & 0 & +1 \end{pmatrix}$$

respectively.

- Notice that there is implicit smoothing in these operators, and the magnitudes of the derivatives are scaled up.

## Additional Steps are Needed Following Smoothing and Gradient Computation

- After gradient computation, we still just have an image
- Pure thresholding leaves us with pixel locations of strong gradients, not edges (edge elements).
- We want to "extract" the edges, both individually and in contours.
- Decisions about what is a true edge and what is not may be made in groups of edgels.
- We'll sketch a picture in class to illustrate this.



## Non-Maximum Suppression

- For each pixel, suppress it as a "non-edge" if a neighboring pixel has a stronger gradient.
- We'll start with a 1-d image where the gradients are shown in the following (with the column index above the boxes and the gradient magnitude in the boxes):

*gradient mag*  
*x-section*

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
4	6	7	8	11	15	18	17	8	6	4	8	10	8	4

and demonstrate the importance of non-maximum suppression and not recording the values "in-place" in the array.

*must compare to  
"suppressed possible edge bcs*

*if threshold is 8 we get two edges/  
peaks*

*"thinning"*



- In 2d it is important to think about what should remain after non-maximum suppression.
- We will focus our discussion on the following example where the gradient magnitudes are shown as integers (assume blank boxes have gradient 0):

	0	1	2	3	4	5	6	7
0						85	94	71
1					78	100	96	
2				55	82	98	80	
3			49	80	102	99	65	
4			71	100	104	71		
5		72	90	89	68			
6	73	91	92	76				
7	80	99	80	45				

gradient mag image!



get blue

8-connected  
get 3 peaks  
4-connected  
get 4

- Discussion:

- What happens if we don't consider gradient direction during non-maximum suppression?
- How to use gradient direction? add orange
- Discretization effects must be considered here, since the determination of who is a "neighbor" may not be symmetric.

- Finally, candidate edgels are the pixels that are left after non-maximum suppression.

- This can include locations with very low gradient magnitudes!

trick to use grad dir to do NMs → grad dir to choose one of 4 pairs of direction  $-1,-1$  and  $1,1$



$0,-1$

$1,-1$

$0,1$

$-1,1$

$-1,0$   $1,0$

- look at 2 nbrs along dir (one each way)
- "suppression" grad (as non max) if one of these nbrs has larger gradient

## Sub-Pixel Localization of Edges

Still only know pixel locations of peaks, so...

- At each surviving peak, fit a parabola to the gradient magnitudes along the gradient direction.
- Offset the edge location along the normal based on the location of the peak.
- Yields the subpixel edge location
- While fitting a parabola sounds like an expensive computation, we will see in class that it reduces to straightforward arithmetic.
- Per usual we will start with 1d and then consider 2d.

## Edge Linking

Group edgels into chains:

- Tangent direction is  $90^\circ$  rotated from the gradient direction.
- Record “next” and “previous” edgels at each edge (think of a doubly-linked list) — the edges “ahead” and “behind” the current one.
- Resolving ambiguities:
  - Find candidates for linking based on consistency of position **and** orientation.
  - If there are two candidates ahead (or behind) and they are consistent with each other, choose the four-connected neighbor first
  - If they aren’t consistent, then this is a branching location. We can either
    - \* Terminate the chain or
    - \* Attempt to find the stronger link and continue it.
- Result is a series of edgel chains.

## Edge Thresholding and Hysteresis

We still haven't decided which points are really "edgels", so we need to apply a threshold:

- Could use a single global threshold, perhaps computed from the image gradient statistics.
  - For example, assume the edgels occupy no more than a certain percentage of the image and use the histogram to compute a threshold.
- Could compute different thresholds in different (but overlapping) regions, perhaps interpolating between them to generate (potentially) different thresholds at each pixel.
- Double thresholding — also called "hysteresis" thresholding:
  - Two thresholds:  $\theta_1 < \theta_2$ .
  - All edgels with gradient magnitude below  $\theta_1$  are eliminated
  - Those above  $\theta_2$  are kept.
  - Those in between must be connected by a chain of edgels with strengths above  $\theta_1$  to an edgel with strength above  $\theta_2$  in order to survive.

## Examples

Using Python and OpenCV, we will examine edge detection results on a number of images and consider how well it works.

## Summary: How Well Does it Work?

- When thinking about how we want a vision system to work, we think “boundaries” instead of “edges”.
- Therefore, we want “edges” where there are no significant intensities changes, and we want to ignore many prominent edges.
- Some of this can be seen in the following examples from Martin, Fowlkes, Malik, IEEE T-PAMI 2004:




Figure 1: How well does edge detection work?